

Notebook

June 17, 2025

```
[6]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime, timedelta
from io import StringIO
import warnings
warnings.filterwarnings('ignore')

# Time series libraries
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from sklearn.metrics import mean_absolute_error, mean_squared_error
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

# Set style for better plots
try:
    plt.style.use('seaborn-v0_8')
except:
    plt.style.use('seaborn')
sns.set_palette("husl")

class ElectricityBillForecaster:
    def __init__(self, csv_data):
        self.df = self.load_and_clean_data(csv_data)
        self.ts_data = None
        self.models = {}
        self.forecasts = {}

    def load_and_clean_data(self, csv_data):
        """Load and clean the electricity bill data"""
        # Read CSV data
        df = pd.read_csv(StringIO(csv_data))

        # Clean column names
```

```

df.columns = df.columns.str.strip()

# Convert dates
df['Bill Date'] = pd.to_datetime(df['Bill Date'])
df['Collection Date'] = pd.to_datetime(df['Collection Date'])

# Clean amount column
df['Transaction Amount'] = pd.to_numeric(df['Transaction Amount'],
errors='coerce')

# Remove any rows with missing amounts
df = df.dropna(subset=['Transaction Amount'])

# Sort by collection date (this represents the billing period end)
df = df.sort_values('Collection Date')

return df

def create_time_series(self):
    """Create time series with proper indexing"""
    # Use Collection Date as index since it represents when bill is paid
    # Each bill represents a 2-month period ending on the collection date
    self.ts_data = pd.Series(
        data=self.df['Transaction Amount'].values,
        index=pd.to_datetime(self.df['Collection Date']),
        name='Bill_Amount'
    )

    print(f"Time series created with {len(self.ts_data)} data points")
    print(f>Date range: {self.ts_data.index.min()} to {self.ts_data.index.
max()})

    return self.ts_data

def analyze_stationarity(self):
    """Check if the time series is stationary"""
    result = adfuller(self.ts_data.dropna())
    print('\nStationarity Test (Augmented Dickey-Fuller):')
    print(f'ADF Statistic: {result[0]:.6f}')
    print(f'p-value: {result[1]:.6f}')

    if result[1] <= 0.05:
        print("Series is stationary")
        return True
    else:
        print("Series is non-stationary")
        return False

```

```

def plot_time_series(self):
    """Plot the original time series and its components"""
    fig, axes = plt.subplots(2, 2, figsize=(15, 10))

    # Original series
    axes[0,0].plot(self.ts_data.index, self.ts_data.values, marker='o',
↳linewidth=2)
    axes[0,0].set_title('Electricity Bill Amount Over Time', fontsize=14,
↳fontweight='bold')
    axes[0,0].set_ylabel('Amount ( )')
    axes[0,0].grid(True, alpha=0.3)

    # Moving averages
    ma_3 = self.ts_data.rolling(window=3).mean()
    ma_6 = self.ts_data.rolling(window=6).mean()

    axes[0,1].plot(self.ts_data.index, self.ts_data.values,
↳label='Original', alpha=0.7)
    axes[0,1].plot(ma_3.index, ma_3.values, label='3-period MA',
↳linewidth=2)
    axes[0,1].plot(ma_6.index, ma_6.values, label='6-period MA',
↳linewidth=2)
    axes[0,1].set_title('Moving Averages')
    axes[0,1].legend()
    axes[0,1].grid(True, alpha=0.3)

    # ACF and PACF
    plot_acf(self.ts_data.dropna(), ax=axes[1,0], lags=10)
    axes[1,0].set_title('Autocorrelation Function')

    plot_pacf(self.ts_data.dropna(), ax=axes[1,1], lags=10)
    axes[1,1].set_title('Partial Autocorrelation Function')

    plt.tight_layout()
    plt.show()

def seasonal_decomposition(self):
    """Perform seasonal decomposition"""
    try:
        # Resample to regular frequency if needed
        ts_regular = self.ts_data.resample('2M').mean().interpolate()

        if len(ts_regular) >= 8: # Need at least 2 complete cycles
            decomposition = seasonal_decompose(ts_regular,
↳model='additive', period=6)

```

```

fig, axes = plt.subplots(4, 1, figsize=(15, 12))

decomposition.observed.plot(ax=axes[0], title='Original')
decomposition.trend.plot(ax=axes[1], title='Trend')
decomposition.seasonal.plot(ax=axes[2], title='Seasonal')
decomposition.resid.plot(ax=axes[3], title='Residual')

plt.tight_layout()
plt.show()

    return decomposition
else:
    print("Not enough data for seasonal decomposition")
    return None
except Exception as e:
    print(f"Seasonal decomposition failed: {e}")
    return None

def fit_arima_model(self):
    """Fit ARIMA model with automatic parameter selection"""
    print("\n=== ARIMA Model ===")

    try:
        # Try different ARIMA parameters
        best_aic = float('inf')
        best_params = None
        best_model = None

        # Parameter ranges to try
        p_range = range(0, 3)
        d_range = range(0, 2)
        q_range = range(0, 3)

        for p in p_range:
            for d in d_range:
                for q in q_range:
                    try:
                        model = ARIMA(self.ts_data, order=(p, d, q))
                        fitted_model = model.fit()

                        if fitted_model.aic < best_aic:
                            best_aic = fitted_model.aic
                            best_params = (p, d, q)
                            best_model = fitted_model
                    except:
                        continue
    
```

```

if best_model:
    print(f"Best ARIMA parameters: {best_params}")
    print(f"AIC: {best_aic:.2f}")

    # Forecast next period
    try:
        forecast_result = best_model.get_forecast(steps=1)
        forecast = forecast_result.predicted_mean.iloc[0]

        # Get confidence intervals
        conf_int = forecast_result.conf_int()
        lower_ci = conf_int.iloc[0, 0]
        upper_ci = conf_int.iloc[0, 1]
    except Exception as e:
        print(f"Using alternative forecast method: {e}")
        forecast = best_model.forecast(steps=1)
        if hasattr(forecast, 'iloc'):
            forecast = forecast.iloc[0]
        elif hasattr(forecast, '__getitem__'):
            forecast = forecast[0] if len(forecast) > 0 else 0
        else:
            forecast = float(forecast)

        lower_ci = forecast * 0.9
        upper_ci = forecast * 1.1

    self.models['ARIMA'] = best_model
    self.forecasts['ARIMA'] = {
        'forecast': forecast,
        'lower_ci': lower_ci,
        'upper_ci': upper_ci
    }

    print(f"ARIMA Forecast: {forecast:.2f}")
    return best_model
else:
    print("Could not fit ARIMA model")
    return None

except Exception as e:
    print(f"ARIMA model failed: {e}")
    return None

def fit_exponential_smoothing(self):
    """Fit Exponential Smoothing model"""

```

```

print("\n=== Exponential Smoothing ===")

try:
    # Try Holt-Winters with different configurations
    models_to_try = [
        {'seasonal': None, 'trend': 'add'},
        {'seasonal': None, 'trend': None},
        {'seasonal': 'add', 'trend': 'add', 'seasonal_periods': 6},
    ]

    best_aic = float('inf')
    best_model = None

    for config in models_to_try:
        try:
            model = ExponentialSmoothing(self.ts_data, **config)
            fitted_model = model.fit()

            if fitted_model.aic < best_aic:
                best_aic = fitted_model.aic
                best_model = fitted_model
        except:
            continue

    if best_model:
        try:
            forecast_result = best_model.forecast(steps=1)

            # Handle different types of forecast results
            if hasattr(forecast_result, 'iloc'):
                forecast = forecast_result.iloc[0]
            elif hasattr(forecast_result, '__getitem__') and
↪len(forecast_result) > 0:
                forecast = forecast_result[0]
            else:
                forecast = float(forecast_result)

            self.models['ExpSmoothing'] = best_model
            self.forecasts['ExpSmoothing'] = {
                'forecast': forecast,
                'lower_ci': forecast * 0.85, # Approximate confidence
↪interval
                'upper_ci': forecast * 1.15
            }

            print(f"Exponential Smoothing Forecast: {forecast:.2f}")
            return best_model

```

```

        except Exception as e:
            print(f"Error in exponential smoothing forecast: {e}")
            return None
    else:
        print("Could not fit Exponential Smoothing model")
        return None

except Exception as e:
    print(f"Exponential Smoothing failed: {e}")
    return None

def simple_forecasting_methods(self):
    """Simple forecasting methods as baseline"""
    print("\n=== Simple Methods ===")

    recent_data = self.ts_data.tail(6).values  # Last 6 bills

    # Simple Moving Average
    sma_forecast = np.mean(recent_data[-3:])  # 3-period moving average

    # Weighted Moving Average (more weight to recent)
    weights = np.array([0.5, 0.3, 0.2])
    wma_forecast = np.average(recent_data[-3:], weights=weights)

    # Linear Trend
    x = np.arange(len(recent_data))
    coeffs = np.polyfit(x, recent_data, 1)
    trend_forecast = coeffs[0] * len(recent_data) + coeffs[1]

    # Seasonal adjustment (summer months typically higher)
    current_month = datetime.now().month
    if current_month in [4, 5, 6, 7, 8, 9]:  # Summer months
        seasonal_factor = 1.1
    else:
        seasonal_factor = 0.9

    seasonal_adjusted = np.mean(recent_data) * seasonal_factor

    self.forecasts['SMA'] = {'forecast': sma_forecast}
    self.forecasts['WMA'] = {'forecast': wma_forecast}
    self.forecasts['Trend'] = {'forecast': trend_forecast}
    self.forecasts['Seasonal'] = {'forecast': seasonal_adjusted}

    print(f"Simple Moving Average: {sma_forecast:.2f}")
    print(f"Weighted Moving Average: {wma_forecast:.2f}")
    print(f"Linear Trend: {trend_forecast:.2f}")

```

```

print(f"Seasonal Adjusted: {seasonal_adjusted:.2f}")

def ensemble_forecast(self):
    """Create ensemble forecast from multiple methods"""
    print("\n=== Ensemble Forecast ===")

    forecasts = []
    weights = []

    for method, result in self.forecasts.items():
        if 'forecast' in result:
            forecasts.append(result['forecast'])
            # Give more weight to advanced methods
            if method in ['ARIMA', 'ExpSmoothing']:
                weights.append(0.3)
            else:
                weights.append(0.1)

    if forecasts:
        # Normalize weights
        weights = np.array(weights) / np.sum(weights)
        ensemble_forecast = np.average(forecasts, weights=weights)

        # Calculate confidence interval
        forecast_std = np.std(forecasts)
        lower_ci = ensemble_forecast - 1.96 * forecast_std
        upper_ci = ensemble_forecast + 1.96 * forecast_std

        self.forecasts['Ensemble'] = {
            'forecast': ensemble_forecast,
            'lower_ci': lower_ci,
            'upper_ci': upper_ci
        }

        print(f"Ensemble Forecast: {ensemble_forecast:.2f}")
        print(f"95% Confidence Interval: {lower_ci:.2f} - {upper_ci:.2f}")

        return ensemble_forecast

    return None

def plot_forecasts(self):
    """Plot historical data with forecasts"""
    fig, ax = plt.subplots(figsize=(15, 8))

    # Plot historical data
    ax.plot(self.ts_data.index, self.ts_data.values,

```



```

        marker='o', linewidth=2, label='Historical Bills', color='blue')

    # Get next forecast date (assuming 2-month billing cycle)
    last_date = self.ts_data.index[-1]
    next_date = last_date + timedelta(days=60) # Approximate 2-month cycle

    # Plot forecasts
    colors = ['red', 'green', 'orange', 'purple', 'brown']
    for i, (method, result) in enumerate(self.forecasts.items()):
        if 'forecast' in result:
            forecast_val = result['forecast']
            ax.scatter(next_date, forecast_val,
                       s=100, color=colors[i % len(colors)],
                       label=f'{method}: {forecast_val:.0f}', zorder=5)

            # Add confidence interval if available
            if 'lower_ci' in result and 'upper_ci' in result:
                ax.errorbar(next_date, forecast_val,
                           yerr=[[forecast_val - result['lower_ci']],
                                  [result['upper_ci'] - forecast_val]],
                           color=colors[i % len(colors)], alpha=0.3,
↪ capsize=5)

    ax.set_title('Electricity Bill Forecasting - June-July 2025',
                 fontsize=16, fontweight='bold')
    ax.set_ylabel('Bill Amount ( )', fontsize=12)
    ax.set_xlabel('Date', fontsize=12)
    ax.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
    ax.grid(True, alpha=0.3)

    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    def summary_report(self):
        """Generate summary report"""
        print("\n" + "="*60)
        print("ELECTRICITY BILL FORECASTING SUMMARY REPORT")
        print("="*60)

        # Basic statistics
        print(f"\nHistorical Data Summary:")
        print(f"Total Bills Analyzed: {len(self.ts_data)}")
        print(f>Date Range: {self.ts_data.index.min().strftime('%Y-%m-%d')} to ↪
↪ {self.ts_data.index.max().strftime('%Y-%m-%d')}")
        print(f"Average Bill: {self.ts_data.mean():.2f}")
        print(f"Median Bill: {self.ts_data.median():.2f}")

```

```

print(f"Highest Bill: {self.ts_data.max():.2f}")
print(f"Lowest Bill: {self.ts_data.min():.2f}")
print(f"Standard Deviation: {self.ts_data.std():.2f}")

# Recent trend
recent_6_avg = self.ts_data.tail(6).mean()
older_6_avg = self.ts_data.iloc[-12:-6].mean() if len(self.ts_data) >= 12
else self.ts_data.head(6).mean()
trend_change = ((recent_6_avg - older_6_avg) / older_6_avg) * 100

print(f"\nRecent Trend Analysis:")
print(f"Recent 6 bills average: {recent_6_avg:.2f}")
print(f"Previous 6 bills average: {older_6_avg:.2f}")
print(f"Trend change: {trend_change:+.1f}%")

# Forecasts
print(f"\nForecasts for June-July 2025 Bill:")
print("-" * 40)
for method, result in self.forecasts.items():
    if 'forecast' in result:
        forecast_val = result['forecast']
        print(f"{method:15s}: {forecast_val:8.2f}")
        if 'lower_ci' in result and 'upper_ci' in result:
            print(f"{'':15s} ( {result['lower_ci']:.2f} - {result['upper_ci']:.2f} )")

# Recommendation
if 'Ensemble' in self.forecasts:
    ensemble_val = self.forecasts['Ensemble']['forecast']
    print(f"\n RECOMMENDED FORECAST: {ensemble_val:.2f}")

# Budget planning
buffer = ensemble_val * 0.15 # 15% buffer
budget_amount = ensemble_val + buffer
print(f" Suggested Budget (with 15% buffer): {budget_amount:.2f}")

print("\n" + "="*60)

def main():
    # Your CSV data
    csv_data = """Sr.No.,Reciept no,Bill Date,Collection Date,Transaction
    ↳Amount,Payment Type,Mode of Payment
    1,BD015153BAHAAAPI4MY8,2025-05-02,2025-06-02,2082.68,Energy Charges,INSTA PAY
    2,BD015091BAGAAANTVLDB,2025-03-01,2025-04-01,1404.38,Energy Charges,INSTA PAY
    3,BD015029BAGAAAL12CJI,2024-12-29,2025-01-29,609.54,Energy Charges,INSTA PAY
    4,BD014337BAFAAAKD8VZB,2024-11-02,2024-12-02,2087.56,Energy Charges,INSTA PAY
    5,QF477134,2024-09-03,2024-10-03,1773.0,Energy Charges,CASH

```

```

6,BD014213BAEAAAF3X0QK,2024-06-30,2024-07-31,2425.36,Energy Charges,INSTA PAY
7,BD014151BADAAAD8FB4E,2024-04-30,2024-05-30,2027.82,Energy Charges,INSTA PAY
8,20757185371,2024-03-04,2024-04-04,1348.35,Energy Charges,INSTA PAY
9,20406018170,2024-01-01,2024-02-01,1156.51,Energy Charges,INSTA PAY
10,20094073250,2023-11-01,2023-12-01,2136.14,Energy Charges,INSTA PAY
11,19781106927,2023-08-29,2023-09-29,2084.58,Energy Charges,INSTA PAY
12,19488250821,2023-07-02,2023-08-02,2318.0,Energy Charges,INSTA PAY
13,18944656409,2023-04-29,2023-05-29,2248.0,Energy Charges,INSTA PAY
14,18404042407,2023-02-27,2023-03-27,1483.5,Energy Charges,INSTA PAY
15,15261870624,2022-12-26,2023-01-26,1413.1,Energy Charges,INSTA PAY
16,13095313465,2022-10-28,2022-11-28,2053.34,Energy Charges,INSTA PAY
17,11807371536,2022-09-01,2022-10-01,2182.6,Energy Charges,INSTA PAY
18,11171590790,2022-06-28,2022-07-28,2040.0,Energy Charges,INSTA PAY
19,10958307048,2022-05-01,2022-06-01,1632.62,Energy Charges,INSTA PAY
20,10762870341,2022-03-04,2022-04-04,1064.31,Energy Charges,INSTA PAY
21,9168851987,2022-01-04,2022-02-04,1088.05,Energy Charges,INSTA PAY
22,1652060412,2021-11-06,2021-12-06,1572.0,Energy Charges,CASH
23,1936174327,2021-09-07,2021-10-07,1963.0,Energy Charges,CASH
24,QY194192,2021-06-30,2021-07-30,2190.0,Energy Charges,CASH
25,QY132014,2021-05-01,2021-06-01,1700.0,Energy Charges,CASH
26,QY013744,2021-03-08,2021-04-08,1250.0,Energy Charges,CASH
27,QY043270,2021-01-11,2021-02-11,810.0,Energy Charges,CASH
28,QG496991,2020-11-14,2020-12-14,1490.0,Energy Charges,CASH
29,QG490559,2020-09-08,2020-10-08,2170.0,Energy Charges,CASH
30,SE134053,2020-07-07,2020-08-07,1450.0,Energy Charges,CASH
31,SE130525,2020-05-09,2020-06-09,2120.0,Energy Charges,CASH""

```

```

# Create forecaster instance
forecaster = ElectricityBillForecaster(csv_data)

# Create time series
forecaster.create_time_series()

# Analyze data
forecaster.analyze_stationarity()

# Plot time series
forecaster.plot_time_series()

# Seasonal decomposition
forecaster.seasonal_decomposition()

# Fit models and generate forecasts
forecaster.fit_arima_model()
forecaster.fit_exponential_smoothing()
forecaster.simple_forecasting_methods()

```

```

# Create ensemble forecast
forecaster.ensemble_forecast()

# Plot all forecasts
forecaster.plot_forecasts()

# Generate summary report
forecaster.summary_report()

if __name__ == "__main__":
    main()

```

Time series created with 31 data points


Date range: 2020-06-09 00:00:00 to 2025-06-02 00:00:00


Stationarity Test (Augmented Dickey-Fuller):

ADF Statistic: -1.580460

p-value: 0.493401

Series is non-stationary





=== ARIMA Model ===

Best ARIMA parameters: (0, 1, 2)

AIC: 460.47

ARIMA Forecast: 1830.41

=== Exponential Smoothing ===

Exponential Smoothing Forecast: 2284.22

=== Simple Methods ===

Simple Moving Average: 1365.53

Weighted Moving Average: 1142.62

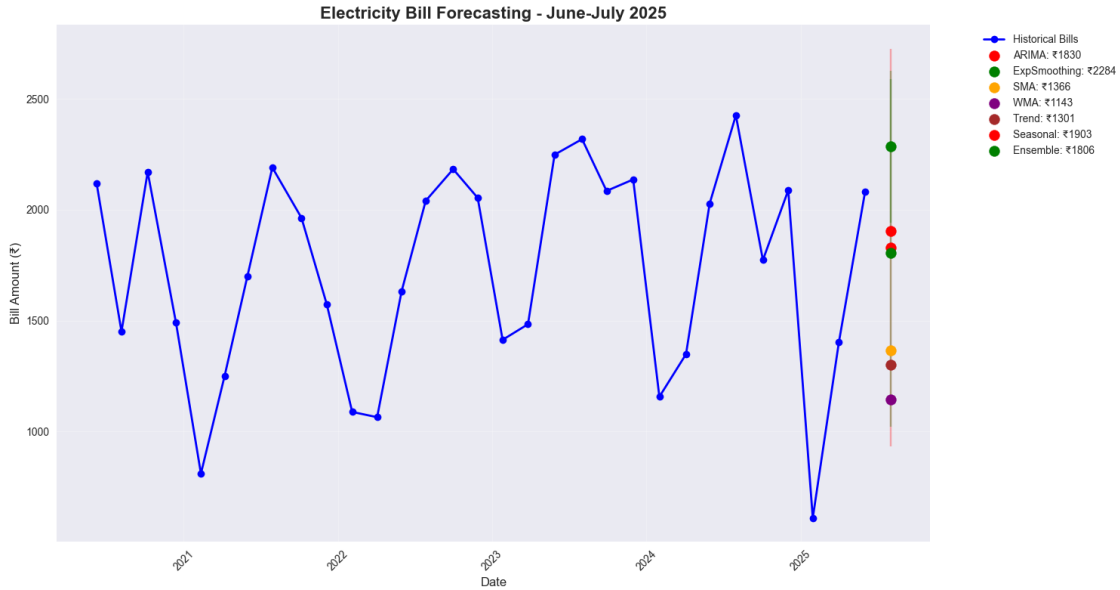
Linear Trend: 1300.69

Seasonal Adjusted: 1903.46

=== Ensemble Forecast ===

Ensemble Forecast: 1805.62

95% Confidence Interval: 1022.23 - 2589.01



ELECTRICITY BILL FORECASTING SUMMARY REPORT

Historical Data Summary:

Total Bills Analyzed: 31
 Date Range: 2020-06-09 to 2025-06-02
 Average Bill: 1721.76
 Median Bill: 1773.00
 Highest Bill: 2425.36
 Lowest Bill: 609.54
 Standard Deviation: 477.72

Recent Trend Analysis:

Recent 6 bills average: 1730.42
 Previous 6 bills average: 1845.23
 Trend change: -6.2%

Forecasts for June-July 2025 Bill:

ARIMA	:	1830.41
		(934.24 - 2726.59)
ExpSmoothing	:	2284.22
		(1941.58 - 2626.85)
SMA	:	1365.53
WMA	:	1142.62
Trend	:	1300.69

Seasonal : 1903.46
Ensemble : 1805.62
(1022.23 - 2589.01)

RECOMMENDED FORECAST: 1805.62
Suggested Budget (with 15% buffer): 2076.46

=====

[]:

This notebook was converted with convert.ploomber.io